

WPF Verteilte Systeme

A distributed Architecture for training a Deep Learning Agent

Jonas Faßbender

Amjad ...

Hamza ...

David ...

Contents

1	Introduction	4
2	Technologies	5
2.1	Python	5
2.1.1	History	5
2.1.2	What is Python?	6
2.1.3	Features and philosophy	7
2.1.4	Why we took Python for our project	8
2.2	Neural Nets	9
2.2.1	Derivation from biology	9
2.2.2	Structure of an artificial neuron	9
2.2.3	Construction of a neural network	11
2.2.4	Learning algorithms	12
2.2.5	Reinforcement Learning	12
2.2.6	Neural Networks and Deep Reinforcement Learning	13
2.3	OpenAI Gym	16
2.4	Tensorflow	17
2.5	Message Broker	18
2.5.1	History	18
2.5.2	Functionality and architecture	18
2.6	RabbitMQ	19
2.6.1	Small glance on the history of RabbitMQ	20
3	Application	21
3.1	The agent	21
3.2	Architecture	23
3.2.1	Network	23
3.2.2	Executioner	24
3.2.3	Worker	24
3.2.4	Queues	24
3.3	Results	24
3.4	Where to go	24

List of Figures

1	Python Logo	5
2	Popular software written in Python	7
3	Scheme of an artificial neuron	9
4	An artificial neuron with bias	10
5	Neural network with hidden layer	11
6	Convolutional Agent[2]	13
7	Difference between a Deep Learning Network and a Neural Network	14
8	Reinforcement learning diagram[1]	15
9	Timeline of message queuing	18
10	Scheme of a message queue	19
11	RabbitMQ Logo	19
12	Architecture of our Application on the network level	23

1 Introduction

This document records our attempt of building a distributed application for training a Deep Learning Agent to master a specific environment.

It should be noted that the focus of this project lies on "distributed", which means the main goal of this application is speeding up a task which requires lots of computing power using technologies and algorithms for concurrent computing in order to horizontally scale up and reduce the overall time the task takes for completing rather than an attempt to find a optimized algorithm that reduces time complexity of training a Deep Learning Agent.

We will begin with introducing the technologies we used, starting with a summary of the Python programming language we have used for building our application, followed by a brief introduction to Neural Networks. After that we will outline the frameworks and libraries we added in order to build our Machine Learning Infrastructure, before finally presenting the concept of a Message Broker, which is one crucial part of our application's architecture.

After that we will present the architecture and the concepts of our application with emphasis on how we achieve concurrency on different levels. This is followed by a chapter which outlines our test results. The document is concluded by giving an overview over improvements and optimizations which can further increase the results of our application.

The whole project can be found at:

<https://github.com/jofas/wpfvs>

2 Technologies

2.1 Python



Figure 1: Python Logo

At the point when computers came to presence in their initial years, there was the need for them to be modified.

Computer scientists came with a 0 and 1 programming language as the optimum solution at that time, yet it was extremely troublesome and tedious. This trouble caused great advancements in programming and prompted the processes of making high level languages that we know and hear about nowadays.

Python is a general-purpose high-level programming language[24] whose design philosophy emphasizes code readability.[25]

Python aims to combine "remarkable power with very clear syntax"[26] and its standard library is large and comprehensive. Its use of indentation for block delimiters is unusual among popular programming languages.

Python can be considered as a multi-paradigm programming language. Rather than forcing programmers to adopt a particular style of programming, it permits several styles: object-oriented programming and structured programming also are fully supported. Many other paradigms are supported using extensions, such as pyDBC[27] and Contracts for Python[28] which allow Design by Contract.

Python uses dynamic typing and a combination of reference counting and a cycle-detecting garbage collector for memory management. An important feature of Python is dynamic name resolution (late binding), which binds method and variable names during program execution.

Rather than requiring all desired functionality to be built into the language's core, Python was designed to be highly extensible. New built-in modules can be easily written in C, C++. Python can also be used as an extension language for existing modules and applications that need a programmable interface. This design of a small core language with a large standard library and an easily extensible interpreter was intended by Van Rossum from the very start because of his frustrations with ABC (which espoused the opposite mindset).[29]

Like other dynamic languages, Python is often used as a scripting language, but is also used in a wide range of non-scripting contexts. Using third-party tools, such as Py2exe or Pyinstaller[30], Python code can be packaged into standalone executable programs. Python interpreters are available for many operating systems.

2.1.1 History

Python was conceived in the late 1980s[3] and its implementation was getting started in December 1989 [4] by Guido van Rossum at CWI in the Netherlands as a successor to the ABC programming language (itself inspired by SETL)[5] capable of exception handling and interfacing with the Amoeba operating system. [6]

Van Rossum is Python's principal author, and his continuing central role in deciding the direction of Python is reflected in the title given to him by the Python community, Benevolent Dictator for Life (BDFL).

Python 2.0 was released on 16 October 2000, with many major new features including a full garbage collector and support for Unicode. However, the most important change was to the development process itself, with a shift to a more transparent and community-backed process.

Python 3.0, a major, backwards-incompatible release, was released on 3rd of December 2008^[7] after a long period of testing. Many of its major features have been back-ported to the backwards-compatible Python 2.7.^[8]

2.1.2 What is Python?

Python is a high-level programming language designed to be easy to read and simple to implement. It is open source, which means it is free to use, even for commercial applications.

Python can run on Mac, Windows, and UNIX operating systems and has also been ported to Java and .NET virtual machines. ^[9]

Python is considered a scripting language, like Ruby or Perl and is often used for creating Web applications and dynamic Web content. It is also supported by a number of 2D and 3D imaging programs, enabling users to create custom plug-ins and extensions with Python. Examples of applications that support a Python API include GIMP, Blender, and Autodesk Maya.^[9]

Python supports the use of modules and packages, which means that programs can be designed in a modular style and code can be reused across a variety of projects. Once the user develops a module or package she/he needs, it can be scaled for use in other projects, and it's easy to import or export these modules.^[10]

Scripts written in Python (.py files) can be parsed and run immediately. They can also be saved as a compiled programs (.pyc files), which are often used as programming modules that can be referenced by other Python programs. ^[9]

That being said, it can be considered that python has a lot of distinguishing properties and features as being portable since, it can be used on a lot of operating systems. In addition there are some other properties such:

1. Python supports other technologies:

It can support COM, .Net, etc. objects. Also, some alternatives and complements were created for Python that make it easier to work with these objects in an integrated mode.

2. Presence of Third Party Modules:

The Python Package Index (PyPI) contains numerous third-party modules that make Python capable of interacting with most of the other languages and platforms.^[11]

3. Python is open source:

Even though all rights of this program are reserved for the Python institute, but it is open source and there is no limitation in using, changing and distributing.^[12]

4. Learning Ease and Support Available:

Python offers excellent readability and uncluttered simple-to-learn syntax which helps beginners to utilize this programming language. The code style guidelines, PEP 8, provide a set of rules to facilitate the formatting of code. Additionally, the wide base of users and active developers has resulted in a rich internet resource bank to encourage development and the continued adoption of the language.^[11]

5. Productivity and Speed:

Python has clean object-oriented design, provides enhanced process control capabilities, and possesses strong integration and text processing capabilities and its own unit testing framework, all of which contribute to the increase in its speed and productivity. Python is considered a viable option for building complex multi-protocol network applications.[11]

2.1.3 Features and philosophy

Python uses dynamic typing and a combination of reference counting and a cycle-detecting garbage collector for memory management. An important feature of Python is dynamic name resolution (late binding), which binds method and variable names during program execution.[14]

The design of Python offers only limited support for functional programming in the Lisp tradition. The language has map, reduce and filter functions, comprehensions for lists, dictionaries, and sets, as well as generator expressions.[34] The standard library has two modules (itertools and functools) that implement functional tools borrowed from Haskell and Standard ML.[15]

An important goal of the Python developers is making Python fun to use. This is reflected in the origin of the name which comes from Monty Python, and in an occasionally playful approach to tutorials and reference materials, for example using spam and eggs instead of the standard foo and bar.[35]

Python has a large standard library, commonly cited as one of Python's greatest strengths, providing tools suited for many tasks. This is deliberate and has been described as a "batteries included" Python philosophy. For Internet-facing applications, a large number of standard formats and protocols (such as MIME and HTTP) are supported. Modules for creating graphical user interfaces, connecting to relational databases, pseudorandom number generators, arithmetic with arbitrary precision decimals, manipulating regular expressions, and doing unit testing are also included.[35]

Some parts of the standard library are covered by specifications (for example, the WSGI implementation wsgiref follows PEP 333), but the majority of the modules are not. They are specified by their code, internal documentation, and test suite (if supplied). However, because most of the standard library is cross-platform Python code, there are only a few modules that must be altered or completely rewritten by alternative implementations.[35]

The standard library is not essential to run Python or embed Python within an application. Blender 2.49 for instance omits most of the standard library.

10 Popular Software Programs Written in Python



Figure 2: Popular software written in Python

2.1.4 Why we took Python for our project

Python is a well-designed language that can be used for real world programming. Some of its strengths were very appealing to us when we thought about which language to use.

The most obvious strength and the reason we decided to use Python are the libraries, frameworks and APIs for machine learning written for/in Python (especially Tensorflow and Keras, which we used, but including far more well designed and beloved tools like PyTorch, Theano, matplotlib, etc).

Here are some other strengths of Python we really appreciated during development:

- System programming:

Internal interfaces of python that are created for working with services of operating system cause Python to be a suitable language for system programming.

These interfaces provide some functions such as: files and directories operations, parallel processing, etc. the standard library of Python can support the different types of platforms and operating systems. It contains some tools for working with system resources such as: environmental variables, files, sockets, pipes, processes, multiple treats, command line, standard stream interfaces, shell programming, etc. [\[35\]](#)

- Network and internet programming:

Various modules are embedded in Python standard library that provide many tools for network programmers, such as: client-server connection, socket programming, FTP, Telnet, email functions, RPC, SOAP, etc.

Also, some third-party tools like mod-Python allow web servers like apache to run Python scripts. Furthermore, some popular programs such as: Django, Turbo gears, Pylon, Zope and Web Ware support Python scripts.[\[35\]](#)

- Numerical programming:

Python with the NumPy library for Linear Algebra computations (used in Tensorflow) can be a powerful alternative for FORTRAN and C++, because it provides powerful tools for working with mathematical libraries, by using simple Python code. Also, there are many third-party tools on the internet for numerical computations.

2.2 Neural Nets

2.2.1 Derivation from biology

”Neural networks are an attempt to model the insights gained in brain research on the interaction between nerve cells (neurons) and connections (synapses)”.[37]

Here, an artificial neuron mimics the functioning of a nerve cell . A biological nerve cell is connected to other nerve cells via synapses. At these connections, stimuli are transmitted by means of chemical messengers (neurotransmitters) and converted into an electrical signal within the nerve cell.

The synapses of a nerve cell are located at so-called dendrites (branches), from which the transmitted stimuli are transmitted to the soma (cell body). It does matter how far the synapse is from the soma. The closer a synapse is to the soma, the stronger the stimulus transmission.[33, 37]

In addition, it should be the case that multiple synapses stimuli at the same time which adds incoming stimuli within the nerve cell.[33] If the stimulus transmitted exceeds a threshold within the cell, then an action potential is triggered and the cell forwards the signal to other cells, which in turn are connected via synapses to the stimulating transmitting cell. Nerve cells form associations in the human brain, which develop the ability to recognize complex structures. Artificial neural networks are an attempt to imitate these associations of nerve cells.[37]

2.2.2 Structure of an artificial neuron




Figure 3: Scheme of an artificial neuron

From chapter 2.2.1, the components and functioning of an artificial neuron can be derived.

In an artificial neuron the following happens:

1. An artificial neuron is connected to other neurons or the direct value input via input connections, which are supposed to represent the synapses of the nerve cell (described in Figure 3 with x_i). These inputs can be discrete or continuous.[36]

If the data comes from other neurons, the weighting of the connection is additionally transferred to the value ω_i . Each connection of a neuron has an individual weighting. The greater the value of this weighting, the more important the transmitted value x_i for the network output.

2. The value and weight of each input connection is added by the transfer function $\sum_{i=1}^n w_i x_i$ to calculate the net input value t .
3. Every artificial neuron, like a nerve cell, has a threshold. An artificial neuron is a value θ . The threshold θ is subtracted from the net input value t to determine the activation potential of the neuron.

The threshold can also be represented by the bias. The bias is a further input, which transmits the constant value 1 and as weighting the value $b = -\theta$.




Figure 4: An artificial neuron with bias

4. The determined value of $t - \theta$ is used as an input in the activation function to determine the output of the neuron.

From this process, the basic elements of a neuron can be determined:

- Weighting
- Threshold
- Transfer function
- Activation function

2.2.3 Construction of a neural network




Figure 5: Neural network with hidden layer

Neural networks are built in several levels (layers). Each of these layers consists of a certain number of neurons and connections. Each neuron of a layer is connected to all neurons of the next layer via the already discussed weighted input connections.

- **Input layer:**

The input layer is the first layer of a neural network. It contains the same Number of neurons, such as input and output values in the test and learning data available.

- **Hidden layer:**

A neural network can not contain, one or more hidden layers. Neural networks can solve simple problems without hidden layers, many problems, such as the so-called "XOR problem", can't be solved without at least one of these layers.[\[36\]](#)

The hidden layers consist of neurons, which together represent an internal representation of the input. The weighting to the individual neurons of the hidden layer is determined, for example, by means of back propagation.

- **Output-Layer:**

The output layer is the last layer of a neural network. It delivers the results of the network to the outside world. The number of neurons in the output layer equals the number of expected output values.

2.2.4 Learning algorithms

Each neural network must be trained with training records. These records consist of rows with the input values and the expected output (supervised learning). The training itself is an iterative process in which each input and output passes through rows.[37]

In order for the network to learn how to approximate the output value, the weights and threshold of each neuron must be modified and adjusted.

The quality measure is the quadratic cost function:

$$E = \frac{1}{2} \sum_{i=1}^n (a_i - o_i)^2$$

The network error E consists of the accumulated squared deviations of the individual rows.

These are calculated from the difference between the expected output value a_i and the calculated network output o_i . The sum is multiplied by $\frac{1}{2}$. This factor simplifies the derivation. The aim now is to find the combination of thresholds and weights of the neurons in which the network error E is the lowest.

You could now let the net simply try all possible weights and thresholds until the combination with the smallest error E is found. This is not possible in practice, as there are usually too many parameters that have to be tried out within the network.[37]

For this purpose, learning algorithms are used instead. Learning algorithms distinguish between two categories, supervised and unsupervised learning.

2.2.5 Reinforcement Learning

Reinforcement learning (RL) is an area of machine learning, inspired by behaviorist psychology, concerned with how software agents ought to take actions in an environment so as to maximize some notion of cumulative reward.

The problem, due to its generality, is studied in many other disciplines, such as game theory, control theory, operations research, information theory, simulation-based optimization, multi-agent systems, swarm intelligence, statistics and genetic algorithms. In the operations research and control literature, reinforcement learning is called approximate dynamic programming, or neuro-dynamic programming.

The problems of interest in reinforcement learning have also been studied in the theory of optimal control, which is concerned mostly with the existence and characterization of optimal solutions, and algorithms for their exact computation, and less with learning or approximation, particularly in the absence of a mathematical model of the environment. In economics and game theory, reinforcement learning may be used to explain how equilibrium may arise under bounded rationality.[1]

While neural networks are responsible for recent breakthroughs in problems like computer vision, machine translation and time series prediction – they can also combine with reinforcement learning algorithms to create something astounding like AlphaGo.[13]

Reinforcement learning refers to goal-oriented algorithms, which learn how to attain a complex objective (goal) or maximize along a particular dimension over many steps; for example, maximize the points won in a game over many moves. They can start from a blank slate, and under the right conditions they achieve superhuman performance. Like a child incentivized by spankings and candy, these algorithms are penalized when they make the wrong decisions and rewarded when they make the right ones – this is reinforcement. Reinforcement algorithms that incorporate deep learning can beat world champions at the game of Go as well as human experts playing numerous Atari video games.

While that may sound trivial, it's a vast improvement over their previous accomplishments, and the state of the art is progressing rapidly.[13]

Reinforcement learning solves the difficult problem of correlating immediate actions with the delayed returns they produce. Like humans, reinforcement learning algorithms sometimes have to wait a while to see the fruit of their decisions. They operate in a delayed return environment, where it can be difficult to understand which action leads to which outcome over many time steps.

Reinforcement learning algorithms can be expected to perform better and better in more ambiguous, real-life environments while choosing from an arbitrary number of possible actions, rather than from the limited options of a video game. That is, with time we expect them to be valuable to achieve goals in the real world.[2]

2.2.6 Neural Networks and Deep Reinforcement Learning

Where do neural networks fit in with RL? Neural networks are the agent that learns to map state-action pairs to rewards. Like all neural networks, they use coefficients to approximate the function relating inputs to outputs, and their learning consists to finding the right coefficients, or weights, by iteratively adjusting those weights along gradients that promise less error.

In reinforcement learning, convolutional networks can be used to recognize an agent's state; e.g. the screen that Mario is on, or the terrain before a drone. That is, they perform their typical task of image recognition.

But convolutional networks derive different interpretations from images in reinforcement learning than in supervised learning. In supervised learning, the network applies a label to an image; that is, it matches names to pixels.

In reinforcement learning, given an image that represents a state, a convolutional net can rank the actions possible to perform in that state; for example, it might predict that running right will return 5 points, jumping 7, and running left none.

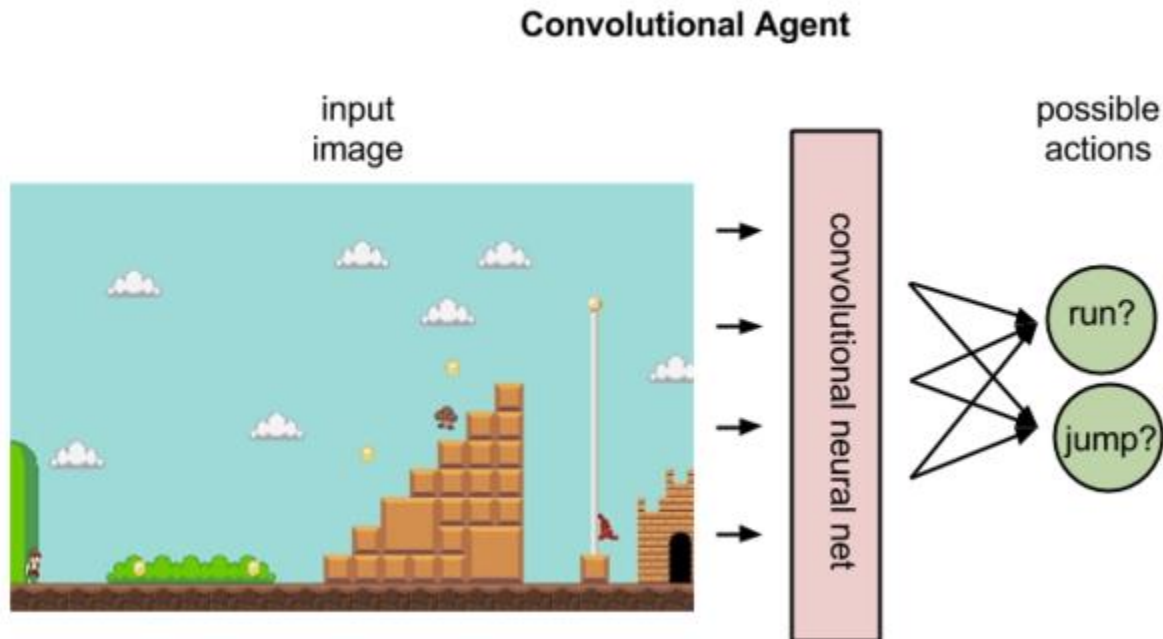



Figure 6: Convolutional Agent[2]

The above image illustrates what a policy agent does, mapping a state to the best action.

At the beginning of reinforcement learning, the neural network coefficients may be initialized stochastically, or randomly. Using feedback from the environment, the neural net can use the difference between its expected reward and the ground-truth reward to adjust its weights and improve its interpretation of state-action pairs.


Reinforcement learning relies on the environment to send it a scalar number in response to each new action. The rewards returned by the environment can be varied, delayed or affected by unknown variables, introducing noise to the feedback loop.[2]

Simple Neural Network



● Input Layer

Deep Learning Neural Network



● Hidden Layer

● Output Layer

Figure 7: Difference between a Deep Learning Network and a Neural Network

The typical framing of a Reinforcement Learning (RL) scenario: an agent takes actions in an environment, which is interpreted into a reward and a representation of the state, which are fed back into the agent.




Figure 8: Reinforcement learning diagram[1]

2.3 OpenAI Gym

”OpenAI Gym is a toolkit for reinforcement learning research. It includes a growing collection of benchmark problems that expose a common interface, and a website where people can share their results and compare the performance of algorithms.”[31]

OpenAI Gym is an easy to use and easy to integrate toolkit with a variety of environments for reinforcement learning (RL) (cmp. 2.2.6) which we used for training our agent.

```
# a small example program using OpenAI Gym

import gym
import random

# build gym environment. A full list of available environ-
# ments can be found at: https://gym.openai.com/envs/
environment = gym.make('LunarLander-v2')

# the range of actions an agent can perform
action_space = environment.action_space.n

# mandatory first reset
environment.reset()

# generate random action (instead of predicted action from
# the agent)
action = random.randrange(0, action_space)

# observations provided by the gym environment after
# doing the random action in the environment
observation, reward, done, info = environment.step(action)
```


2.4 Tensorflow

2.5 Message Broker

A message broker is an intermediary computer program module that translates a message from the formal messaging protocol of the sender to the formal messaging protocol of the receiver.

Message brokers are elements in telecommunication or computer networks where software applications communicate by exchanging formally-defined messages. Message brokers are a building block of message-oriented middleware (MOM) but are typically not a replacement for traditional middleware like MOM and remote procedure call (RPC). [21]

2.5.1 History

Message broker products are the middleware that embody different messaging solutions. In 1983 Vivek Ranadivé from Teknekron Software Systems began working on an idea based on the ideals of a Software Bus to enable applications to share data in a standard fashion. That work would later be referred to as “The Information Bus”. Already in 1986 his ideas were put to use when Goldman Sachs, an American investment banking firm, launched cooperation with Teknekron to find solutions for the trading floor of the future.[22]

For nearly two decades the domain of message exchange was left for proprietary vendors and proprietary formats. Throughout the 1980s and 1990s message queuing kept evolving but in isolation since commercial message queue vendors strived for interoperability between client applications rather than worked on standardized interfaces for their message queuing products to utilize. [22]

The figure bellow shows a brief timeline of message queuing:

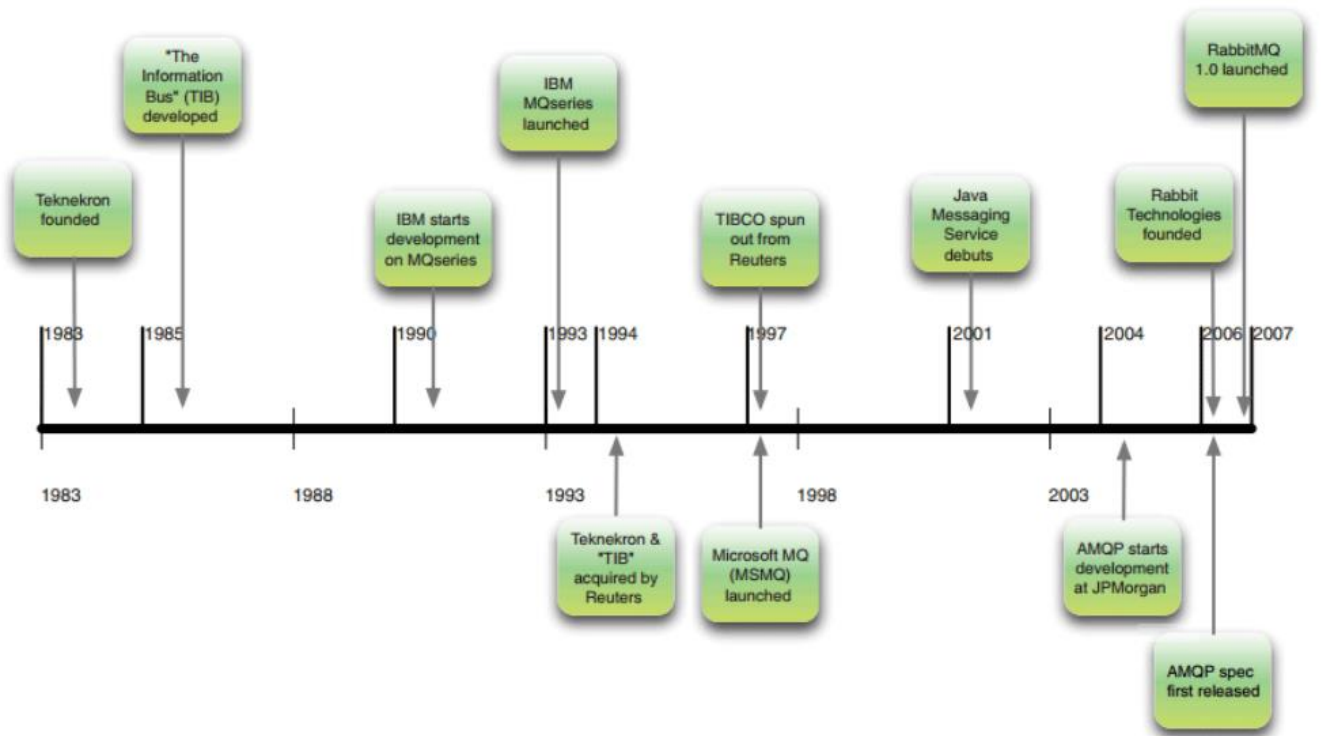


Figure 9: Timeline of message queuing

2.5.2 Functionality and architecture

A message broker is an architectural pattern for message validation, transformation, and routing.

It mediates communication among applications, minimizing the mutual awareness that applications should have of each other in order to be able to exchange messages, effectively implementing decoupling.[32]

The primary purpose of a broker is to take incoming messages from applications and perform some action on them. Message brokers can decouple endpoints, meet specific non- functional requirements, and facilitate reuse of intermediary functions. For example, a message broker may be used to manage a workload queue or message queue for multiple receivers, providing reliable storage, guaranteed message delivery and perhaps transaction management. [21]

Message brokers are generally based on one of two fundamental architectures: hub-and-spoke and message bus. In the prior, a central server acts as the mechanism that provides integration services, whereas with the latter, the message broker is a communication backbone or distributed service that acts on the bus.[23]




Figure 10: Scheme of a message queue

2.6 RabbitMQ



Figure 11: RabbitMQ Logo

For our project we used RabbitMQ as our message broker.

RabbitMQ is an open source message broker software (sometimes called message-oriented middleware) that originally implemented the Advanced Message Queuing Protocol(AMQP) and has since been extended with a plug-in architecture to support Streaming Text Oriented Messaging Protocol (STOMP), Message Queuing Telemetry Transport (MQTT), and other protocols.[16]

The RabbitMQ server program is written in the Erlang programming language and is built on the Open Telecom Platform framework for clustering and failover. Client libraries to interact with the broker are available for all major programming languages.[17]

2.6.1 Small glance on the history of RabbitMQ

Rabbit Technologies Ltd., originally developed RabbitMQ. Rabbit Technologies started as a joint venture between LShift and CohesiveFT in 2007[18], and was acquired in April 2010 by SpringSource, a division of VMware.[19] The project became part of Pivotal Software in May 2013.[20]

3 Application

Our goal was to make use of a lot of computing power in order to train our agent to master an OpenAI Gym Environment (cmp. 2.3). In this chapter we will document how we tried to achieve this goal with our distributed application.

3.1 The agent

We programmed our agent as a neural network with the Keras library, which has an API for high level, high abstraction neural networks.

Keras uses Tensorflow as its backend for computations and basically only provides a nicer abstraction of Tensorflow.

```
# a small example program using Keras
#
# API documentation at: https://keras.io/

# Sequential is the keras object representing a neural net-
# work
from keras.models import Sequential
# a neural net is comprised of layers connected with each
# other. The Dense object represents a layer
from keras.layers import Dense

# the neural net. This specific neural net has four layers
# (the input (size of 12), two hidden (both 64 artificial
# neurons) and the output layer (size of 4)). The input
# layer does not have to be specified since the input_dim
# parameter of the first layer automatically generates the
# input layer.
model = Sequential([
    Dense(64, activation='relu', input_dim=12 ),
    Dense(64, activation='relu' ),
    Dense(4, activation='softmax' ),
])

# define how the model should learn and some other meta
# information for the training process (learning algorithm,
# optimizer, etc)
model.compile(
    optimizer = 'adam',
    loss      = 'categorical_crossentropy',
    metrics   = ['accuracy']
)

# generate a dummy data set with corresponding labels with
# 1000 entries for training
```

```
import numpy as np

data = np.random.random((1000,12))
# generating the labels (either a 0 or a 1)
labels = np.random.randint(2, size=(1000, 1))

# training the model, iterating 10 times
model.train(data, labels, epochs=10)

# using the neural net to predict the label of a random
# data point
test = np.random.random((1,12))

model.predict(test)
```

3.2 Architecture

This chapter will describe in detail how our application is build, how it works and what we use to achieve concurrency.

On a higher level our application is devided into two distinct parts, the Executioner (E) and the Worker (W).

3.2.1 Network

Instances of both, E_i and W_i communicate over a message broker, in this case RabbitMQ (cmp. 2.5, 2.6).




Figure 12: Architecture of our Application on the network level

3.2.2 Executioner

3.2.3 Worker

here comes the data sanitation part

3.2.4 Queues

3.3 Results

3.4 Where to go

(dynamic data generation, loadbalancing, further optimizations like thread pool exec in worker (not spawning so often), destroying the bottleneck in executioner (data parsing))

References

- [1] Wikipedia. https://en.wikipedia.org/wiki/Reinforcement_learning#/media/File:Reinforcement_learning_diagram.svg. 17.07.2018.
- [2] Artificial Intelligence Wiki. <https://skymind.ai/wiki/deep-reinforcement-learning>. 17.07.2018.
- [3] The Making of Python. <http://www.artima.com/intv/pythonP.html>. 20.07.2018.
- [4] A brief Timeline of Python. <http://python-history.blogspot.com/2009/01/brief-timeline-of-python.html>. 20.07.2018.
- [5] van Rossum, Guido. "[Python-Dev] SETL (was: Lukewarm about range literals)". <http://mail.python.org/pipermail/python-dev/2000-August/008881.html>. 20.07.2018.
- [6] Why was Python created in the first place? <http://www.python.org/doc/faq/general/#why-was-python-created-in-the-first-place>. 20.07.2018.
- [7] Python 3.0 Release. <http://python.org/download/releases/3.0/>. 20.07.2018.
- [8] PEP 3000 – Python 3000. <http://www.python.org/dev/peps/pep-3000/>. 20.07.2018.
- [9] Python Definition. <https://techterms.com/definition/python>. 20.07.2018.
- [10] what is python? <http://www.pythonforbeginners.com/learn-python/what-is-python/>. 20.07.2018.
- [11] Benefits of Python. <https://www.invensis.net/blog/it/benefits-of-python-over-other-programming-languages/>. 20.07.2018.
- [12] About python. <https://www.python.org/about/>. 20.07.2018.
- [13] AlphaGo. <https://deepmind.com/research/alphago>. 17.07.2018.
- [14] Python Programming language. [https://en.wikipedia.org/wiki/Python_\(programming_language\)](https://en.wikipedia.org/wiki/Python_(programming_language)). 20.07.2018.
- [15] Itertools. <https://docs.python.org/3/library/itertools.html>. 20.07.2018.
- [16] Which protocols does RabbitMQ support? <https://www.rabbitmq.com/protocols.html>. 20.07.2018.
- [17] RabbitMQ. <https://en.wikipedia.org/wiki/RabbitMQ>. 20.07.2018.
- [18] Launch of RabbitMQ Open Source Enterprise Messaging. http://www.rabbitmq.com/resources/RabbitMQ_PressRelease_080207.pdf. 20.07.2018.
- [19] Rabbit Technologies announce acquisition by SpringSource. <https://web.archive.org/web/20100418132113/http://www.rabbitmq.com/news.html>. 20.07.2018.
- [20] Proudly part of Pivotal. <https://web.archive.org/web/20130602054940/http://www.rabbitmq.com/news.html>. 20.07.2018.
- [21] Message Broker. https://en.wikipedia.org/wiki/Message_broker. 20.07.2018.

- [22] Message brokers and RabbitMQ in Action. https://www.theseus.fi/bitstream/handle/10024/78263/Kamppuri_Tsuri.pdf?sequence=1. 20.07.2018.
- [23] Integration Brokers and Web Services. https://books.google.de/books?id=2EonCgAAQBAJ&pg=PA71&redir_esc=y#v=onepage&q&f=false. 20.07.2018.
- [24] General Python FAQ. <http://docs.python.org/faq/general.html#what-is-python-good-for>. 20.07.2018.
- [25] What is Python? Executive Summary. <http://www.python.org/doc/essays/blurb/>. 20.07.2018.
- [26] General Python FAQ. <http://www.python.org/doc/faq/general/#what-is-python>. 20.07.2018.
- [27] Contracts for Python. <http://www.nongnu.org/pydbc>. 20.07.2018.
- [28] Contracts for Python. <http://www.wayforward.net/pycontract/>. 20.07.2018.
- [29] The Making of Python. <http://www.artima.com/intv/pythonP.html>. 20.07.2018.
- [30] PyInstaller Home Page. <https://www.pyinstaller.org/>. 20.07.2018.
- [31] BROCKMAN, G., CHEUNG, V., PETTERSSON, L., SCHNEIDER, J., SCHULMAN, J., TANG, J., AND ZAREMBA, W. OpenAI Gym, 2016.
- [32] EJSMONT. "Asynchronous Processing". Web Scalability for Startup Engineers. McGraw Hill Professional, 2015.
- [33] ERDMANN, A., ERDMANN, U., AND MARTENS, A. Grüne Reihe. Materialien für den Sekundarbereich II. Schroedel Verlag GmbH, Braunschweig, 2004.
- [34] HOETTINGER, AND RAYMOND. PEP 289 - Generator Expressions. Python Software Foundation, 2002.
- [35] NOSRATI, M. Python: An appropriate language for real world programming.
- [36] REY, G. D., AND BECK, F. Neuronale Netze - Eine Einführung. http://www.neuronaletesnetz.de/downloads/neuronaletesnetz_de.pdf. 17.07.2018.
- [37] TRINKWALDER, A. Netzgespinste. c't 6 (2016), 130–135.